



# CNN with TF



**Rishit Dagli**

10 STD, TEDX, Ted-Ed speaker|Google certified mobile site developer|Intel AI Scholar|2XGCP Champ|Mozilla Mumbai Lead

—  
<http://bit.ly/tfugm>

<http://bit.ly/tfugm-p>



People with no idea about AI,  
Wondering AI will destroy humanity









Me wondering why my NN is classifying a  
cat as a dog





<https://www.kaggle.com/c/dogs-vs-cats/data>

  Competitions Datasets Kernels Discussion Learn ...  



## Dogs vs. Cats

Create an algorithm to distinguish dogs from cats

215 teams · 5 years ago


[Overview](#) [Data](#) [Kernels](#) [Discussion](#) [Leaderboard](#) [Rules](#) [Team](#)

### Data Description


The training archive contains 25,000 images of dogs and cats. Train your algorithm on these files and predict the labels for test1.zip (1 = dog, 0 = cat).

### A note on hand labeling

Per the rules and spirit of this contest, please do not manually label your submissions. We work hard to fair and fun contests, and ask for the same respect in return.









# Activity Recognition



```
if(speed<4){  
    status=WALKING;  
}
```





# Activity Recognition



```
if(speed<4){  
    status=WALKING;  
}
```



```
if(speed<4){  
    status=WALKING;  
} else {  
    status=RUNNING;  
}
```



# Activity Recognition



```
if(speed<4){  
    status=WALKING;  
}
```



```
if(speed<4){  
    status=WALKING;  
} else {  
    status=RUNNING;  
}
```



```
if(speed<4){  
    status=WALKING;  
} else if(speed<12){  
    status=RUNNING;  
} else {  
    status=BIKING;  
}
```



# Activity Recognition



```
if(speed<4){  
    status=WALKING;  
}
```



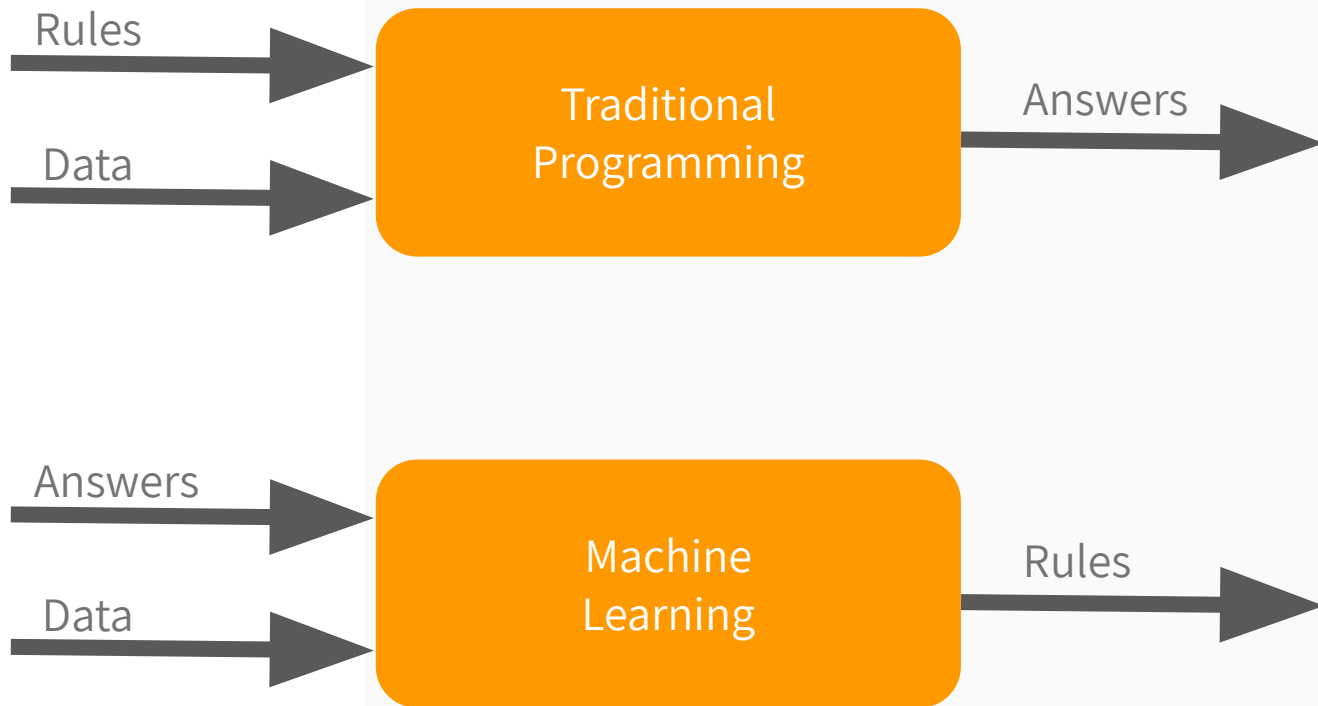
```
if(speed<4){  
    status=WALKING;  
} else {  
    status=RUNNING;  
}
```



```
if(speed<4){  
    status=WALKING;  
} else if(speed<12){  
    status=RUNNING;  
} else {  
    status=BIKING;  
}
```



// Oh crap





# Activity Recognition



```
0101001010100101010
1001010101001011101
0100101010010101001
0101001010100101010
```

Label = WALKING



```
1010100101001010101
0101010010010010001
0010011111010101111
1010100100111101011
```

Label = RUNNING



```
10010100111111010101
1101010111010101110
1010101111010101011
1111110001111010101
```

Label = BIKING



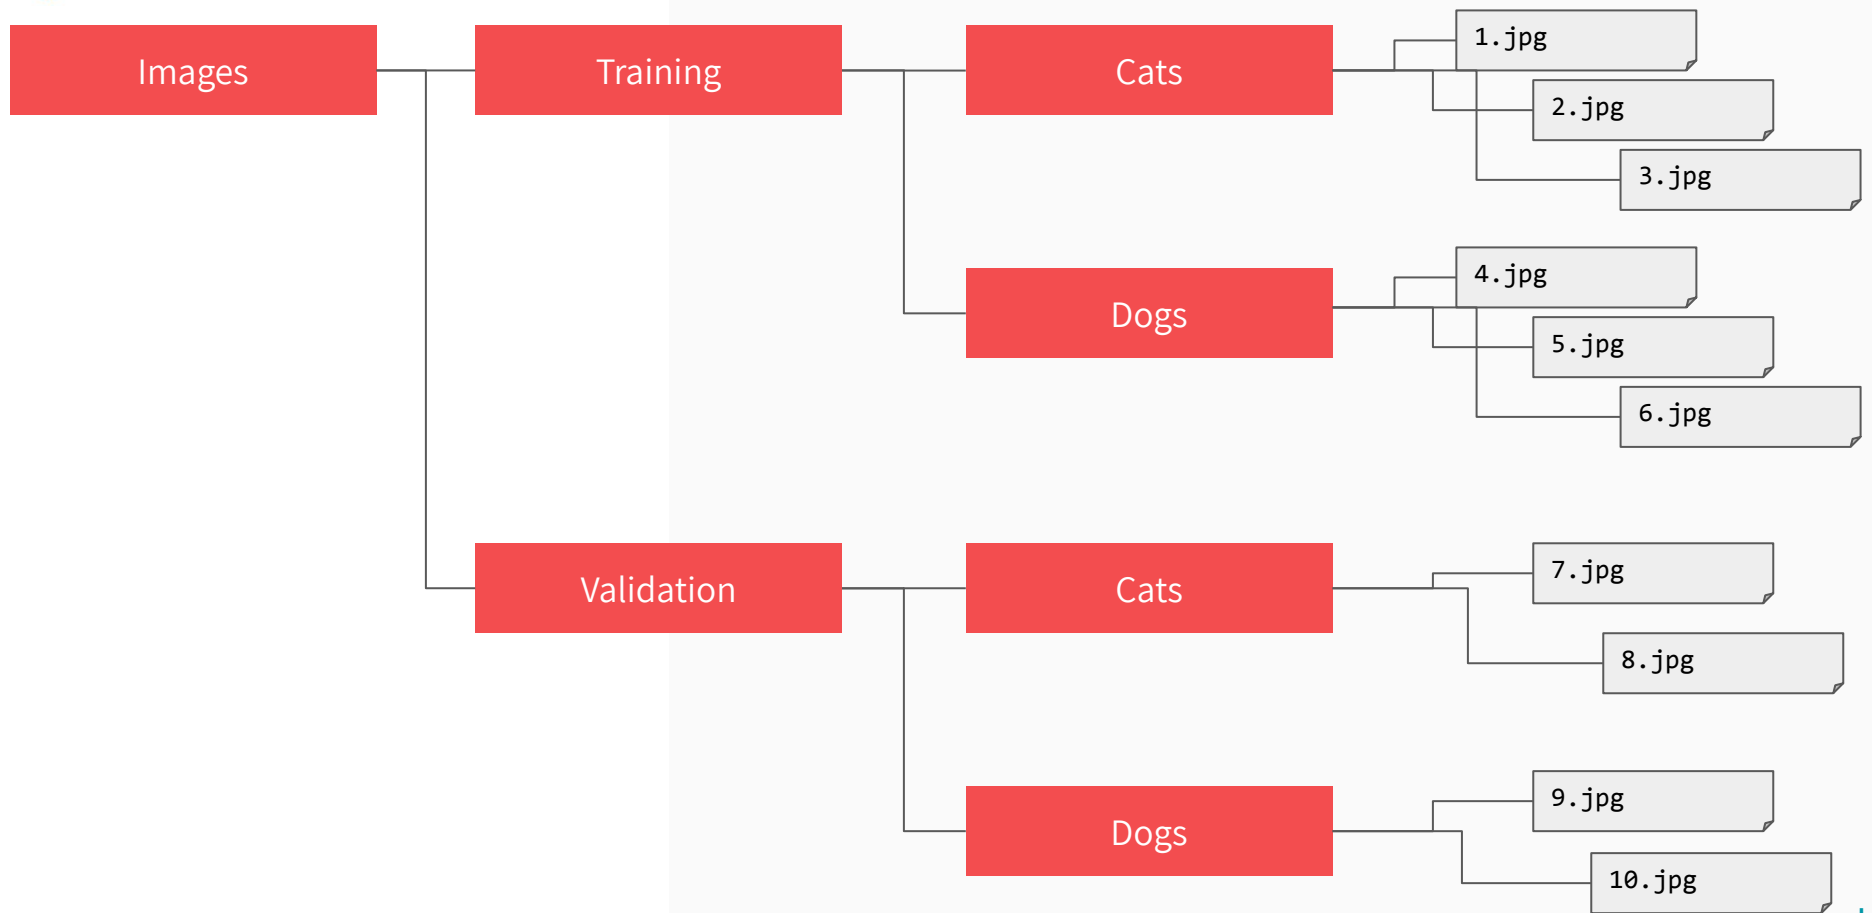
```
11111111111010011101
0011111010111110101
0101110101010101110
1010101010100111110
```

Label = GOLFING  
(Sort of)

# What to expect???

- Building *data input pipelines* using the `tf.keras.preprocessing.image.ImageDataGenerator` class to efficiently work with data on disk to use with the model.
- Build and test model.
- *Overfitting* —How to identify and prevent it.
- *Data augmentation* and *dropout* —techniques to fight overfitting

# Understanding the Data







Images

Training

Cats

1.jpg

2.jpg

3.jpg

Dogs

4.jpg

5.jpg

6.jpg

Validation

Cats

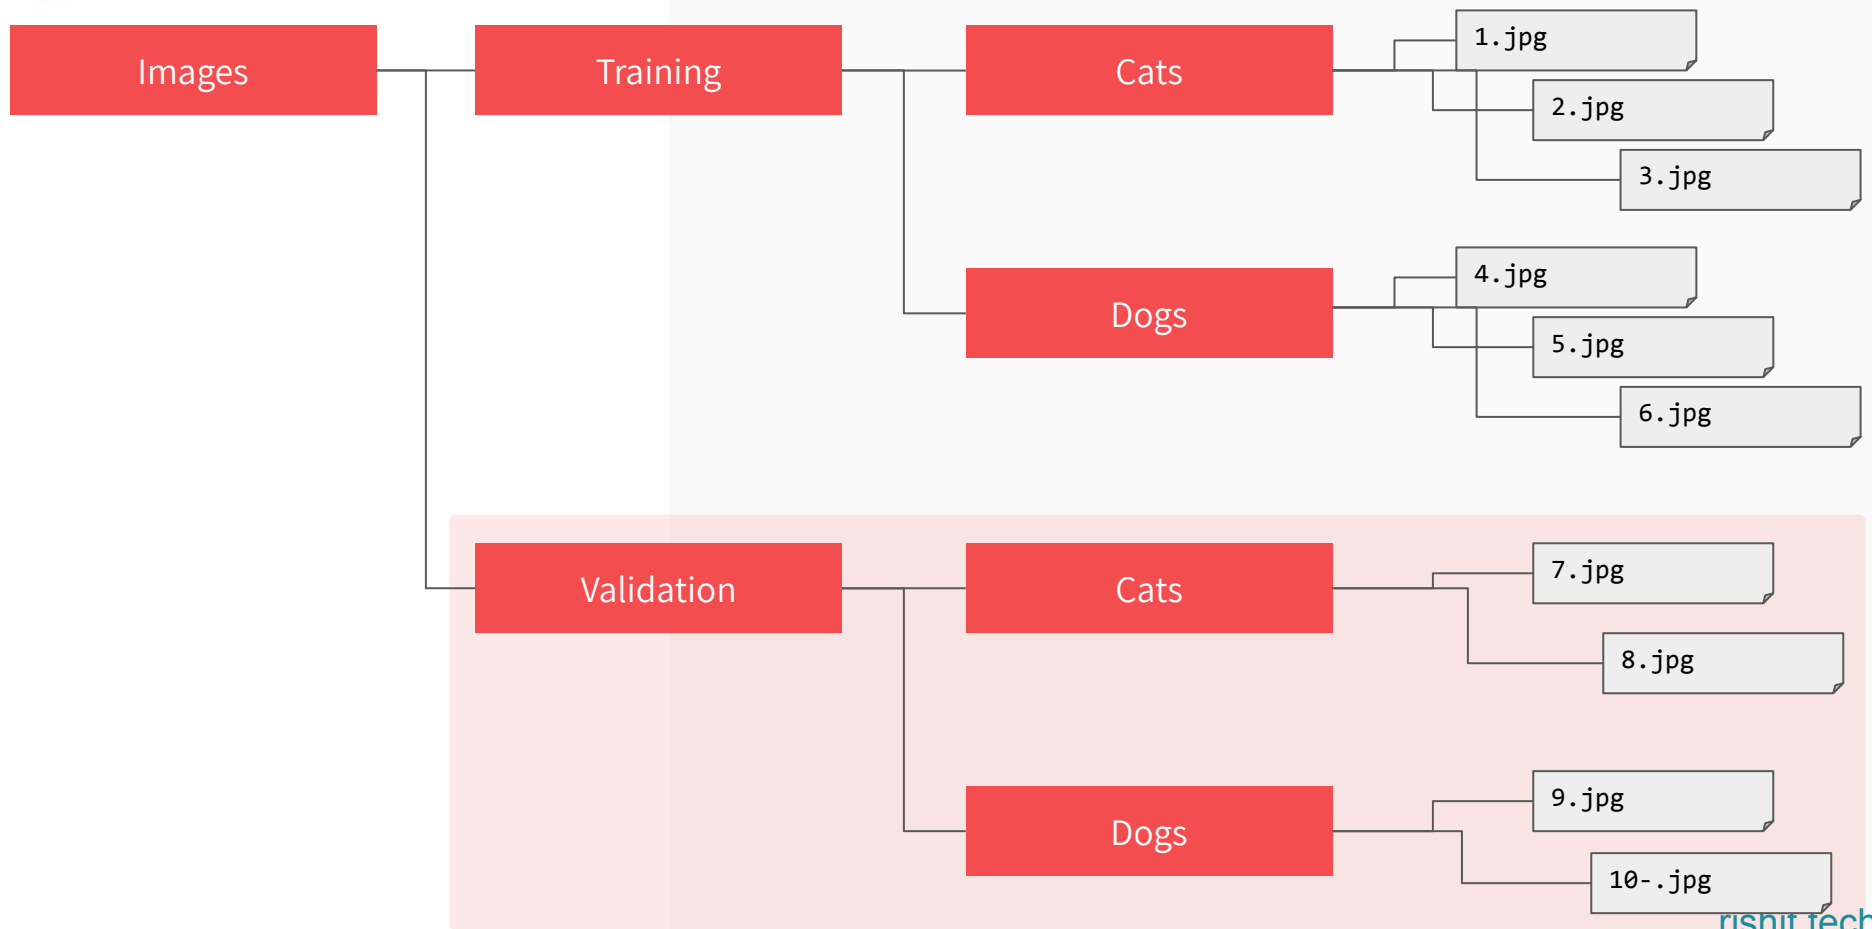
7.jpg

8.jpg

Dogs

9.jpg

10.jpg





```
total training cat images: 1000
total training dog images: 1000
total validation cat images: 500
total validation dog images: 500
--
Total training images: 2000
Total validation images: 1000
```

# Imports



```
import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers
import Dense, Conv2D, Flatten, Dropout, MaxPooling2D
from tensorflow.keras.preprocessing.image
import ImageDataGenerator

import os
import numpy as np
import matplotlib.pyplot as plt
```



# Parameters

`batch_size = 128`

`epochs = 15`

`IMG_HEIGHT = 150`

`IMG_WIDTH = 150`



# Data Preparation

1. Read images from the disk.
2. Decode contents of these images and convert it into proper grid format as per their RGB content.
3. Convert them into floating point tensors.
4. Rescale the tensors from values between 0 and 255 to values between 0 and 1, as neural networks prefer to deal with small input values.



```
from tensorflow.keras.preprocessing.image  
import ImageDataGenerator
```



```
# Generator for our training data
```

```
train_image_generator = ImageDataGenerator(rescale=1./255)
```

```
train_data_gen = train_image_generator.flow_from_directory(  
    batch_size=batch_size,  
    directory=train_dir,  
    shuffle=True,  
    target_size=(IMG_HEIGHT, IMG_WIDTH)  
    class_mode='binary')
```

```
# Generator for our training data
```

```
train_image_generator = ImageDataGenerator(rescale=1./255)
```

```
train_data_gen = train_image_generator.flow_from_directory(  
    batch_size=batch_size,  
    directory=train_dir,  
    shuffle=True,  
    target_size=(IMG_HEIGHT, IMG_WIDTH)  
    class_mode='binary')
```

```
# Generator for our training data
```

```
train_image_generator = ImageDataGenerator(rescale=1./255)
```

```
train_data_gen = train_image_generator.flow_from_directory(  
    batch_size=batch_size,  
    directory=train_dir,  
    shuffle=True,  
    target_size=(IMG_HEIGHT, IMG_WIDTH)  
    class_mode='binary')
```

```
# Generator for our training data
```

```
train_image_generator = ImageDataGenerator(rescale=1./255)
```

```
train_data_gen = train_image_generator.flow_from_directory(  
    batch_size=batch_size,  
    directory= train_dir,  
    shuffle=True,  
    target_size=(IMG_HEIGHT,IMG_WIDTH)  
    class_mode='binary')
```

```
# Generator for our training data
```

```
train_image_generator = ImageDataGenerator(rescale=1./255)
```

```
train_data_gen = train_image_generator.flow_from_directory(
```

```
    batch_size=batch_size,
```

```
    directory= train_dir,
```

```
    shuffle=True,
```

```
    target_size=(IMG_HEIGHT, IMG_WIDTH)
```

```
    class_mode='binary')
```

```
# Generator for our training data
```

```
train_image_generator = ImageDataGenerator(rescale=1./255)
```

```
train_data_gen = train_image_generator.flow_from_directory(  
    batch_size=batch_size,  
    directory= train_dir,  
    shuffle=True,  
    target_size=(IMG_HEIGHT, IMG_WIDTH)  
    class_mode='binary' )
```

```
# Generator for our validation data
```

```
validation_image_generator = ImageDataGenerator(rescale=1./255)
```

```
val_data_gen = validation_imadata_generator.flow_from_directory  
(  
    batch_size=batch_size,  
    directory= validation_dir,  
    shuffle=True,  
    target_size=(IMG_HEIGHT,IMG_WIDTH)  
    class_mode='binary')
```

# Model

Let's understand and Build a Convolutional Neural Network





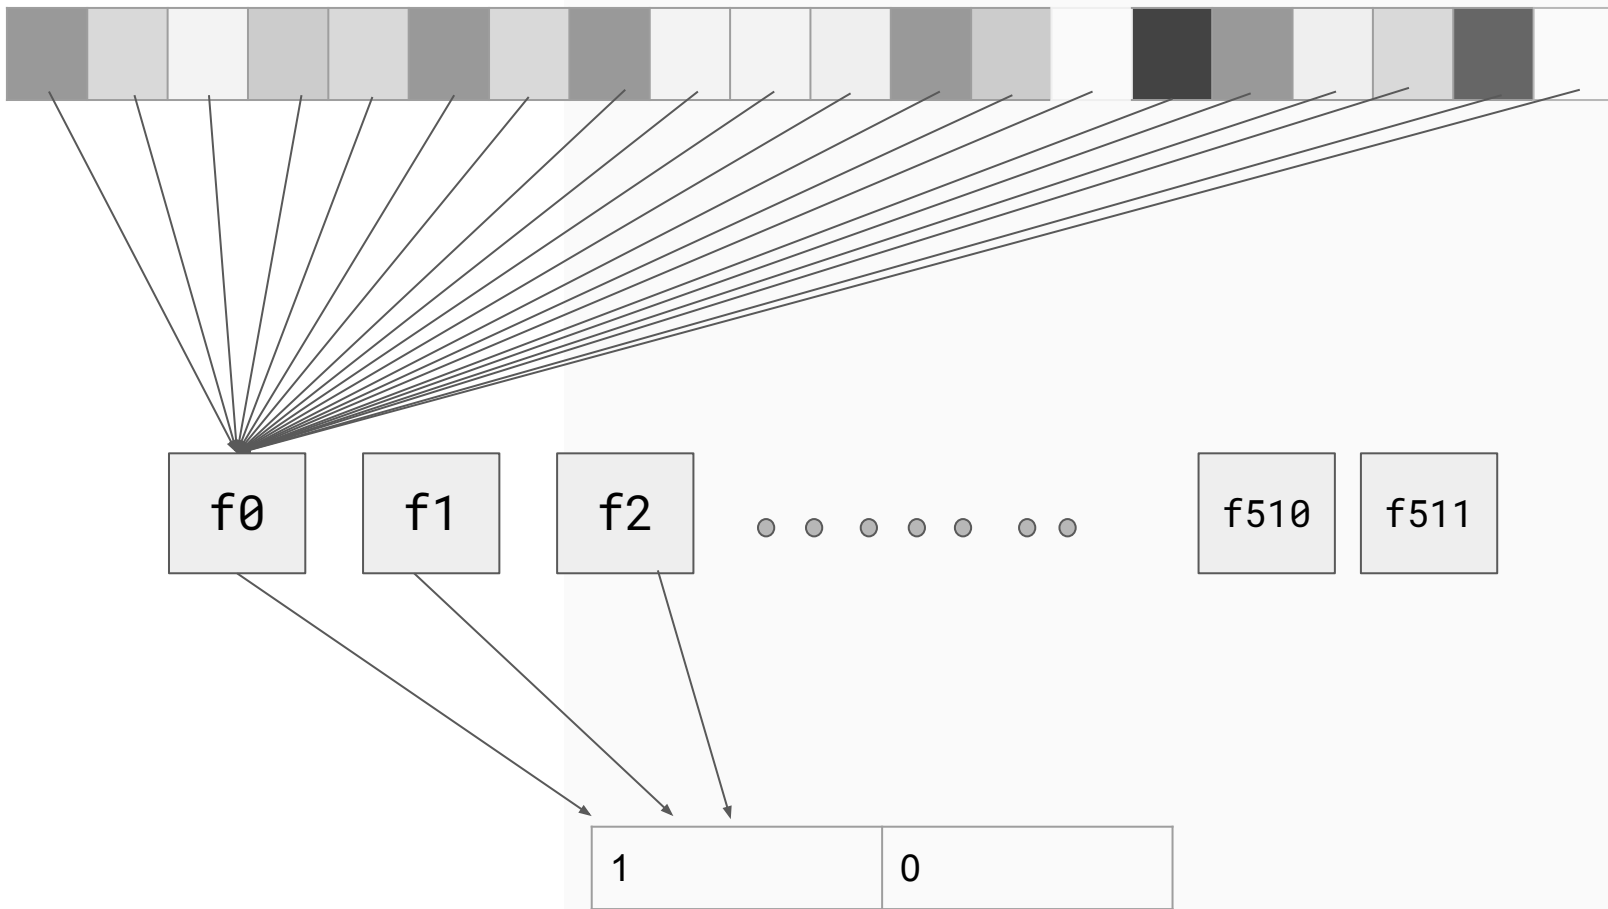
```
model = Sequential([  
    Flatten( input_shape = (150,150,3),  
    Dense(512 , activation = 'relu',  
    Dense(2, activation='softmax')  
])
```



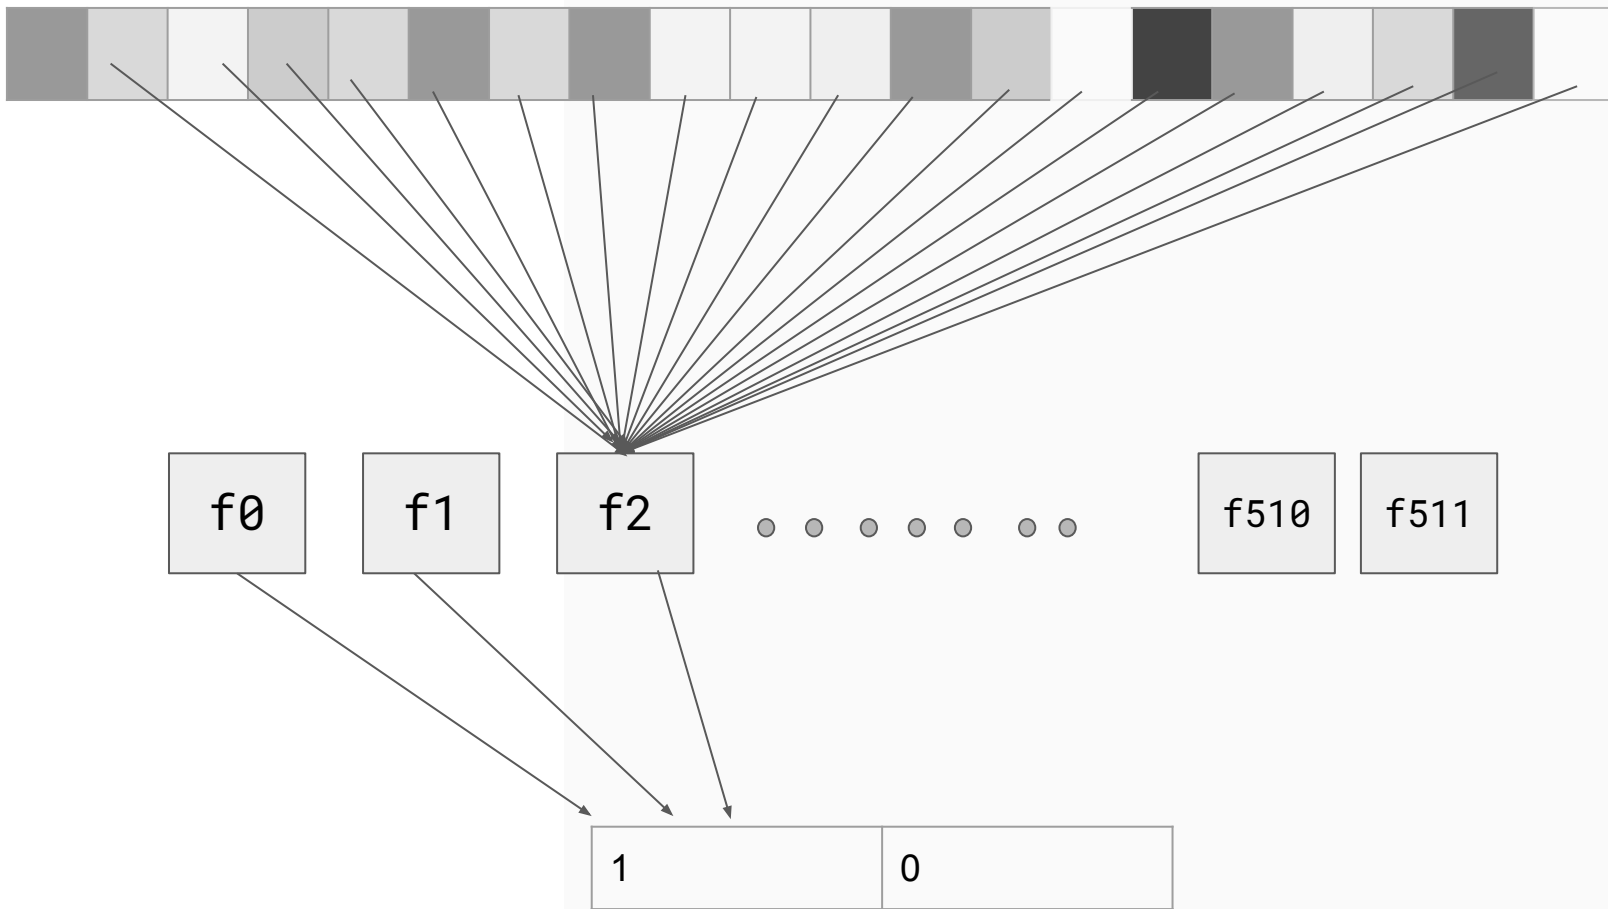
```
model = Sequential([  
    Flatten( input_shape = (150,150,3),  
    Dense(512 , activation = 'relu',  
    Dense( 2 , activation='softmax')  
])
```

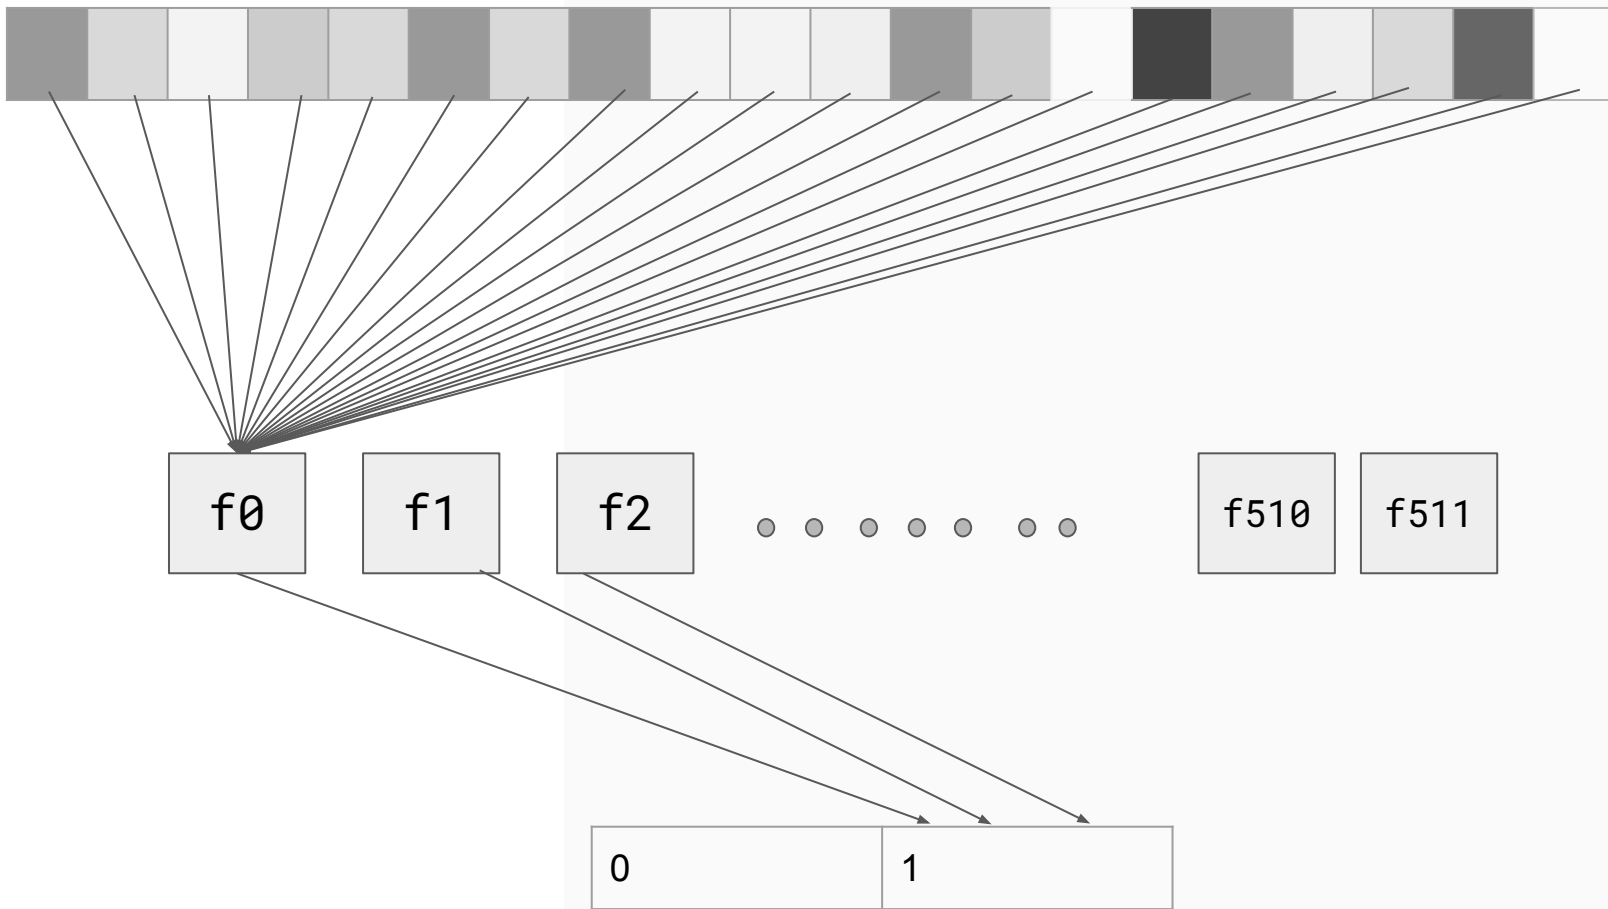


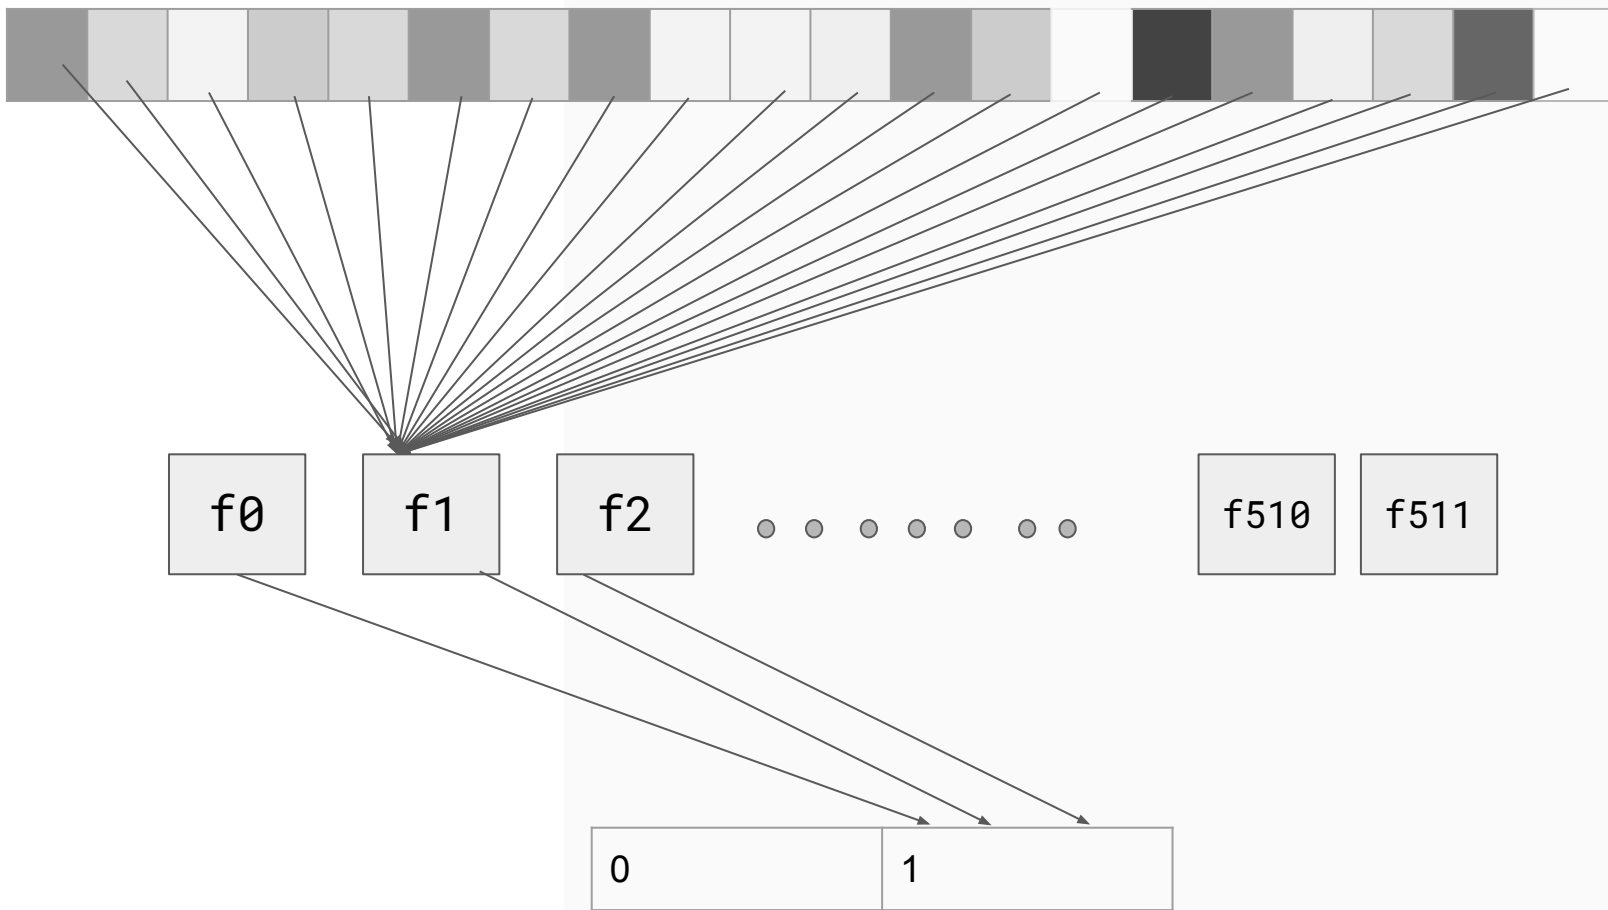
```
model = Sequential([  
    Flatten(input_shape = (150,150,3),  
    Dense(512, activation = 'relu',  
    Dense(2, activation='softmax')  
])
```



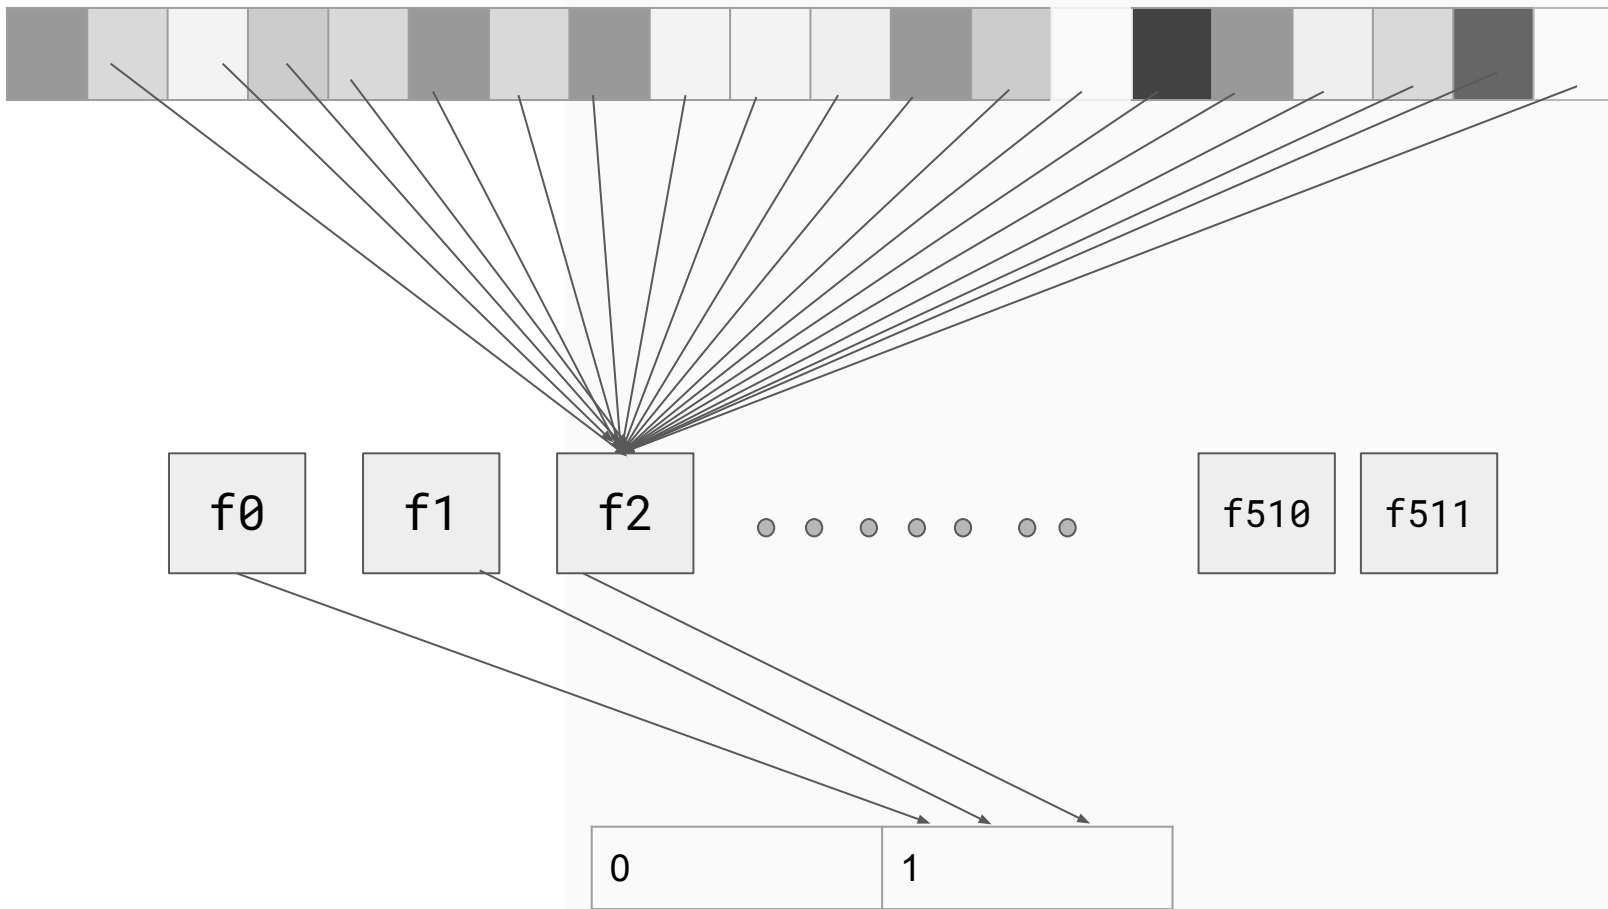













```
model.compile(optimizer='adam',  
              loss='binary_crossentropy',  
              )
```

```
model.fit(....., epochs = 100)
```



0	64	128
48	192	144
142	226	168

Current Pixel Value is 192

Consider neighbor Values

-1	0	-2
.5	4.5	-1.5
1.5	2	-3


Filter Definition

CURRENT\_PIXEL\_VALUE = 192

NEW\_PIXEL\_VALUE =  $(-1 * 0) + (0 * 64) + (-2 * 128) +$   
 $(.5 * 48) + (4.5 * 192) + (-1.5 * 144) +$   
 $(1.5 * 42) + (2 * 226) + (-3 * 168)$




	-1	0	1	
	-2	0	2	
	-1	0	1	





	-1	-2	-1	
	0	0	0	
	1	2	1	





```
model = Sequential([  
    Conv2D(16, (3,3), padding='same', activation='relu',  
    input_shape=(IMG_HEIGHT, IMG_WIDTH ,3)),  
    MaxPooling2D(2,2),  
    Conv2D(32, (3,3), padding='same', activation='relu'),  
    MaxPooling2D(2,2),  
    Conv2D(64, (3,3), padding='same', activation='relu'),  
    MaxPooling2D(2,2),  
    Flatten(),  
    Dense(512, activation='relu'),  
    Dense(1, activation='sigmoid')  
])
```



```
model = Sequential([
    Conv2D(16, (3,3), padding='same', activation='relu',
        input_shape=(IMG_HEIGHT, IMG_WIDTH ,3)),
    MaxPooling2D(2,2),
    Conv2D(32, (3,3), padding='same', activation='relu'),
    MaxPooling2D(2,2),
    Conv2D(64, (3,3), padding='same', activation='relu'),
    MaxPooling2D(2,2),
    Flatten(),
    Dense(512, activation='relu'),
    Dense(1, activation='sigmoid')
])
```



0	64	128	128
48	192	144	144
142	226	168	0
255	0	0	64

0	64
48	192

192

128	128
144	144

144

142	226
255	0

255

168	0
0	64

168

192	144
255	168





# Max Pooling Example



Max Pooling 2X2





```
model = Sequential([
    Conv2D(16, (3,3), padding='same', activation='relu',
        input_shape=(IMG_HEIGHT, IMG_WIDTH ,3)),
    MaxPooling2D(2,2),
    Conv2D(32, (3,3), padding='same', activation='relu'),
    MaxPooling2D(2,2),
    Conv2D(64, (3,3), padding='same', activation='relu'),
    MaxPooling2D(2,2),
    Flatten(),
    Dense(512, activation='relu'),
    Dense(1, activation='sigmoid')
])
```



1	1	0
4	2	1
0	2	1

Pooled Feature Map

Flattening



1
1
0
4
2
1
0
2
1



```
model = Sequential([
    Conv2D(16, (3,3), padding='same', activation='relu',
    input_shape=(IMG_HEIGHT, IMG_WIDTH ,3)),
    MaxPooling2D(2,2),
    Conv2D(32, (3,3), padding='same', activation='relu'),
    MaxPooling2D(2,2),
    Conv2D(64, (3,3), padding='same', activation='relu'),
    MaxPooling2D(2,2),
    Flatten(),
    Dense(512, activation='relu'),
    Dense(1, activation='sigmoid')
])
```



0.3	0.5	0.9	1.0
1.0	1.0	1.0	1.0
0.9	0.9	0.5	0.3
0.2	0.0	0.0	0.0



Input  
 $4 \times 4$




Filter  
 $3 \times 3$




Output  
 $2 \times 2$




```
model = Sequential([
    Conv2D(16, (3,3), padding='same', activation='relu',
    input_shape=(IMG_HEIGHT, IMG_WIDTH ,3)),
    MaxPooling2D(2,2),
    Conv2D(32, (3,3), padding='same', activation='relu'),
    MaxPooling2D(2,2),
    Conv2D(64, (3,3), padding='same', activation='relu'),
    MaxPooling2D(2,2),
    Flatten(),
    Dense(512, activation='relu'),
    Dense(1, activation='sigmoid')
])
```




## Sigmoid

$$\sigma(x) = \frac{1}{1+e^{-x}}$$




## tanh

$$\tanh(x)$$




## ReLU

$$\max(0, x)$$



## Leaky ReLU

$$\max(0.1x, x)$$




## Maxout

$$\max(w_1^T x + b_1, w_2^T x + b_2)$$

## ELU

$$\begin{cases} x & x \geq 0 \\ \alpha(e^x - 1) & x < 0 \end{cases}$$



```
model.compile(optimizer='adam',  
              loss='binary_crossentropy',  
              metrics=['accuracy'])
```






Layer (type)	Output Shape	Param #
conv2d_3 (Conv2D)	(None, 150, 150, 16)	448
max_pooling2d_3 (MaxPooling2D)	(None, 75, 75, 16)	0
conv2d_4 (Conv2D)	(None, 75, 75, 32)	4640
max_pooling2d_4 (MaxPooling2D)	(None, 37, 37, 32)	0
conv2d_5 (Conv2D)	(None, 37, 37, 64)	18496
max_pooling2d_5 (MaxPooling2D)	(None, 18, 18, 64)	0
flatten_1 (Flatten)	(None, 20736)	0
dense_2 (Dense)	(None, 512)	10617344
dense_3 (Dense)	(None, 1)	513
Total params: 10,641,441		
Trainable params: 10,641,441		
Non-trainable params: 0		




```
history = model.fit_generator(  
    train_data_gen,  
    steps_per_epoch=total_train // batch_size,  
    epochs=epochs,  
    validation_data=val_data_gen,  
    validation_steps=total_val // batch_size  
)
```



Training and Validation Accuracy



Training and Validation Loss



# Overfitting



When there are small number of training examples, the model sometimes learns from noises or unwanted details from the training examples- to an extent it negatively impacts the performance of the model on new examples. This is what know as overfitting

How Do You avoid it??

1. Add more data
2. Use data augmentation
3. Use architectures that generalize well.



# Data Augmentation

Apply some random transformation.

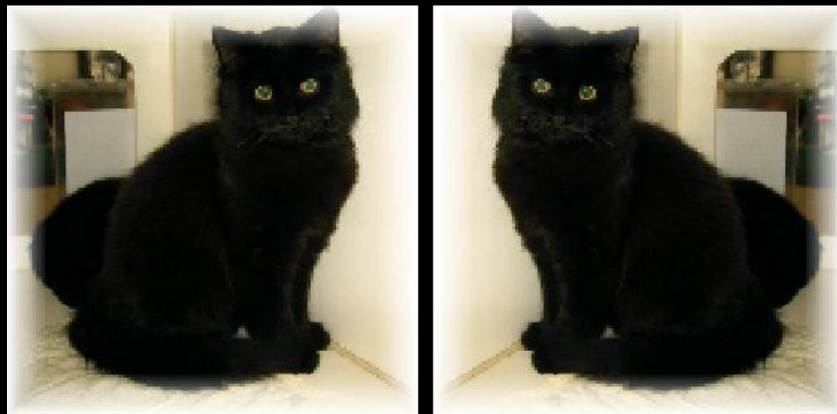
The Goal is the model never sees the exact same picture twice.

Use ImageDataGenerator to perform transformation



# Horizontal Flip

```
image_gen = ImageDataGenerator(rescale=1./255,  
horizontal_flip=True)
```





# Rotate the Image

```
image_gen = ImageDataGenerator(rescale=1./255,  
rotation_range=45)
```







# Zoom Augmentation

```
image_gen = ImageDataGenerator(rescale=1./255,  
                                zoom_range=0.5)
```





## Let's Put this together.

```
image_gen_train = ImageDataGenerator(  
    rescale=1./255,  
    rotation_range=45,  
    width_shift_range=.15,  
    height_shift_range=.15,  
    horizontal_flip=True,  
    zoom_range=0.5  
)
```

```
train_data_gen = image_gen_train.flow_from_directory(  
    batch_size=batch_size,  
    directory=train_dir,  
    shuffle=True,  
    target_size=(IMG_HEIGHT, IMG_WIDTH)  
    class_mode='binary')
```



# DropOut

Another technique to overcome overfitting.

It drops some of the output units from the applied layer during the training process

It takes values such as 0.1, 0.2, 0.3, etc. which means 10%, 20%, 30% of the output dropping off.



# Creating a network with Dropouts

```
model = Sequential([
    Conv2D(16, (3,3), padding='same', activation='relu',
        input_shape=(IMG_HEIGHT, IMG_WIDTH, 3)),
    MaxPooling2D(2,2),
    Dropout(0.2),
    Conv2D(32, (3,3), padding='same', activation='relu'),
    MaxPooling2D(2,2),
    Conv2D(64, (3,3), padding='same', activation='relu'),
    MaxPooling2D(2,2),
    Dropout(0.2),
    Flatten(),
    Dense(512, activation='relu'),
    Dense(1, activation='sigmoid')
])
```



# Testing

```
import numpy as np
from google.colab import files
from keras.preprocessing import image

uploaded = files.upload()


for fn in uploaded.keys():
    path = '/content/' + fn
    img = image.load_img(path, target_size=(150, 150))
    x = image.img_to_array(img)
    x = np.expand_dims(x, axis=0)

    images = np.vstack([x])
    classes = model.predict(images, batch_size=10)
    print(classes[0])
    if classes[0]>0.5:
        print(fn + " is a dog")
    else:
        print(fn + " is a cat")
```



FREE COURSE

# Intro to TensorFlow for Deep Learning

by  **TensorFlow**

This course is a practical approach to deep learning for software developers

START FREE COURSE

NANODEGREE PROGRAM

# Machine Learning Engineer

Become a Machine Learning Engineer

Accelerate your career with the credential that fast-tracks you to job success. >>>

About this Course

COURSE COST

TIMELINE

SKILL LEVEL



Udacity



Explore ▾

What do you want to learn?



For Enterprise

[Log In](#)

Join for Free

**Browse** > **Computer Science** > **Software Development**

This course is part of the **TensorFlow in Practice Specialization**

# Introduction to TensorFlow for Artificial Intelligence, Machine Learning, and Deep Learning

★★★★★ 4.7 4,508 ratings • 899 reviews

## Enroll for Free

Starts Nov 29

### Financial aid available

**100,309** already enrolled!

Offered By



[About](#) [Syllabus](#) [Reviews](#) [Instructors](#) [Enrollment Options](#) [FAQ](#)

11511.0000





**coursera**

Explore ▾

What do you want to learn?



For Enterprise

Log In

Join for Free

Browse > Data Science > Machine Learning

Offered By



# TensorFlow in Practice Specialization

**Enroll for Free**

Starts Nov 29

Financial aid available

**21,914** already enrolled!

[About](#) [How It Works](#) [Courses](#) [Instructors](#) [Enrollment Options](#) [FAQ](#)

## About this Specialization

94,577 recent views

Discover the tools software developers use to build scalable AI-powered algorithms in TensorFlow, a popular open-source machine learning framework.



**100% online courses**

Start instantly and learn at your own schedule